

CSE312

FINAL REPORT

Ahmet Ergani

161044011

DESIGN

Page Table				Virtual Memory		
TLB				Physical Memory		

I created a page table to keep the data about virtual memory entries and a TLB table to keep the data about physical memory entries.

Page Table:

The main purpose of my page table is to directly check whether a certain frame is loaded into memory or not (isValid). And if it is loaded, page table shows the physical memory index of that frame(index).

```
16
17 class PageEntry
18 {
19 public:
20     bool isValid;
21     int index;
22 };
23
```

TLB:

The main purpose of TLB is to check a physical memory space's emptiness (isOccupied) modification status (isModified) and to store which frame is loaded into this space (frameNo). It also

```
31 class TLBEntry
32 {
33 public:
34     bool isOccupied;
35     bool isModified;
36     bool isReferenced;
37     int frameNo;
38     int usage;
39 };
40
```

helps us with the page replacement algorithms. For example I store isReferenced status in it's entries for NRU and a usage integer for LRU

Physical Memory:

I created a dynamic 2d array to hold the frames.

set:

First I calculated which frame has the given index (frameNo). Then I calculated which element of that frame stands for the given index (dataIndex). Then I check the pageTable if this frame is loaded into memory

```
216     if(pageTable[frameNo].isValid)
217     {
218         physicalMemory[pageTable[frameNo].index][dataIndex] = value;
219
220         tlbTable[pageTable[frameNo].index].isModified = true;
221         tlbTable[pageTable[frameNo].index].isReferenced = true;
222         tlbTable[pageTable[frameNo].index].usage = 0;
223     }
```

If it is, I directly set the value of the necessary index of the Physical memory and set the modified bit to true.

Else, a page replacement is accomplished. After that process I set the necessary data in both TLB and Page tables and assign the value to the physical memory index returned to me by replacePage().

```
224     else
225     {
226         physIndex = replacePage(frameNo, tName);
227
228         physicalMemory[physIndex][dataIndex] = value;
229
230         tlbTable[physIndex].isModified = true;
231         tlbTable[physIndex].isOccupied = true;
232         tlbTable[physIndex].frameNo = frameNo;
233         tlbTable[physIndex].usage = 0;
234
235         pageTable[frameNo].isValid = true;
236         pageTable[frameNo].index = physIndex;
237         tlbTable[physIndex].isReferenced = true;
238
239     }
```

get:

This process is nearly same as get. There are only 2 differences between them

- 1) Instead of assigning a value to that index, we just acquire the value of that index and return it
- 2) We don't set the modified bit to true after the operation

replacePage:

I decide which index will be replaced according to the current page replacement algorithm.

```
329         switch(pageReplacement)
330         {
331             case LRU:
332                 physIndex = findLRU();
333                 break;
334
335             case NRU:
336                 physIndex = findNRU();
337                 break;
```

If selected physical memory frame is occupied and has been modified, it's data is saved to disk

```
355         if(tlbTable[physIndex].isOccupied && tlbTable[physIndex].isModified)
356         {
357             lseek(fd, tlbTable[physIndex].frameNo * frameSize * sizeof(int), SEEK_SET);
358
359             for(int i = 0 ; i < frameSize; i++)
360             {
361                 val = write(fd, &physicalMemory[physIndex][i], sizeof(int));
362             }
```

Then the necessary frame is read from disk into selected physical memory index

```
419         lseek(fd, frameNo * frameSize * sizeof(int), SEEK_SET);
420         for(int i = 0 ; i < frameSize; i++)
421         {
422             val = read(fd, &data, sizeof(data));
423             physicalMemory[physIndex][i] = data;
424         }
```

REPLACEMENT ALGORITHMS

LRU:

For this algorithm I stored an integer called usage in TLB. Whenever a set or get occurs every TLB entries' usage increments and the subject entry's usage is set to zero. This way I can know that the entry with the biggest usage value is the least recently used at that moment.

```
428 int Memory::findLRU()
429 {
430     int lruIndex = 0;
431     int max = 0;
432     for(int i = 0; i < physicalMemSize; i++)
433     {
434         if (!tlbTable[i].isOccupied)
435         {
436             return i;
437         }
438         else if(tlbTable[i].usage > max)
439         {
440             lruIndex = i;
441             max = tlbTable[i].usage;
442         }
443     }
444     return lruIndex;
445 }
```

FIFO:

For this algorithm I implemented a basic linked-list queue. When a page replacement occurred I dequeued from this queue and enqueued it afterwards.

```
474 int Memory::findFIFO()
475 {
476     for(int i = 0; i < physicalMemSize; i++)
477     {
478         if (!tlbTable[i].isOccupied)
479         {
480             return i;
481         }
482     }
483     return queue.dequeue();
484 }
```

NRU:

I created a counter and incremented it in set/get operations. I also marked referenced pages during these operations. When counter reaches a certain point these marks are being reset. In case of page replacement I selected the page without referenced mark.

```
447 int Memory::findNRU()
448 {
449     for (int i = 0; i < physicalMemSize; i++)
450     {
451         if (!tlbTable[i].isReferenced)
452         {
453             return i;
454         }
455     }
456     return 0;
457 }
```

SORTING ALGORITHMS

Bubble, quick and merge sort methods are implemented simply and traditionally.

Index sort was different. This sorting algorithm is designed to sort collections of large objects and was useless for our case since we were working with integers. So I acted like my objects were large. Created an indexes array and sorted it with a bubble sort-like

```
656     for (int i = 0; i < N; ++i)
657     {
658         indexes[i] = i + N * 3;
659     }
660
661     for (int i = 0; i < N; i++)
662     {
663         for (int j = i + 1; j < N; j++)
664         {
665             if (get(indexes[i], tName) > get(indexes[j], tName))
666             {
667                 int temp = indexes[i];
668                 indexes[i] = indexes[j];
669                 indexes[j] = temp;
670             }
671         }
672     }
673
674     for (int i = 0; i < N; ++i)
675     {
676         values[i] = get(indexes[i], tName);
677     }
678     for (int i = 0; i < N; ++i)
679     {
680         set(i + N * 3, values[i], tName);
681     }
```

algorithm. Then I created a values array and filled it using the indexes array. Then I set these values to our memory.

THREADS

Each thread sets it's statistics data to zero and calls it's sort method (or directly sorts if the algorithm is not complex).

```
639 void * Memory::mergeSortWorker(void *)
640 {
641     readCount[3] = 0;
642     writeCount[3] = 0;
643     pageReplacementCount[3] = 0;
644     pageMissCount[3] = 0;
645     diskReadCount[3] = 0;
646     diskWriteCount[3] = 0;
647     char tName[6] = "merge";
648     int N = (virtualMemSize * frameSize) / 4;
649     mergeSort(N * 2, N * 3 - 1, tName);
650
651     return NULL;
652 }
```

EXAMPLE RUNS

```
cse312@ubuntu:~/Desktop/final$ ./sortArrays 4 3 5 NRU local 10000 diskFileName.txt
Virtual Memory Size: 32
Physical Memory Size: 8
Frame Size: 16
Replacement Algorithm: NRU
Interval: 10000
File: diskFileName.txt

Array quarters are sorted

FILL : read: 0      write: 512    page miss: 32   page replacement : 24   disk read : 32   disk write : 28
BUBBLE : read: 24126  write: 7870   page miss: 356  page replacement : 348  disk read : 356  disk write : 35
QUICK : read: 1834    write: 918    page miss: 23   page replacement : 15   disk read : 23   disk write : 22
MERGE : read: 896     write: 896    page miss: 27   page replacement : 19   disk read : 27   disk write : 26
INDEX : read: 16384   write: 128    page miss: 104  page replacement : 96   disk read : 104  disk write : 97
CHECK : read: 1016    write: 0      page miss: 24   page replacement : 16   disk read : 24   disk write : 24
```

```
cse312@ubuntu:~/Desktop/final$ ./sortArrays 5 4 6 NRU local 10000 diskFileName.txt
Virtual Memory Size: 64
Physical Memory Size: 16
Frame Size: 32
Replacement Algorithm: NRU
Interval: 10000
File: diskFileName.txt

Array quarters are sorted

FILL : read: 0      write: 2048   page miss: 64   page replacement : 48   disk read : 64   disk write : 56
BUBBLE : read: 395608 write: 133976 page miss: 2734 page replacement : 2718 disk read : 2734   disk write : 2730
QUICK : read: 10829   write: 4984   page miss: 65   page replacement : 49   disk read : 65   disk write : 64
MERGE : read: 4608    write: 4608   page miss: 62   page replacement : 46   disk read : 62   disk write : 61
INDEX : read: 262144 write: 512    page miss: 745  page replacement : 729  disk read : 745  disk write : 740
CHECK : read: 4088    write: 0      page miss: 48   page replacement : 32   disk read : 48   disk write : 48
```

```
cse312@ubuntu:~/Desktop/final$ ./sortArrays 6 5 7 NRU local 10000 diskFileName.txt
Virtual Memory Size: 128
Physical Memory Size: 32
Frame Size: 64
Replacement Algorithm: NRU
Interval: 10000
File: diskFileName.txt

Array quarters are sorted

FILL : read: 0      write: 8192   page miss: 128  page replacement : 96   disk read : 128  disk write : 112
BUBBLE : read: 6222296 write: 2030040 page miss: 27010 page replacement : 26978 disk read : 27010   disk write : 27002
QUICK : read: 59994    write: 17544   page miss: 202  page replacement : 170  disk read : 202  disk write : 201
MERGE : read: 22528    write: 22528   page miss: 198  page replacement : 166  disk read : 198  disk write : 198
INDEX : read: 4194304 write: 2048    page miss: 5023 page replacement : 4991 disk read : 5023   disk write : 4951
CHECK : read: 16376    write: 0      page miss: 96   page replacement : 64   disk read : 96   disk write : 96
```